

Lecture 5: Manipulating Tabular Data

COSC 480 Data Science, Spring 2017
Michael Hay

Goals for today

- Wrap up data processing from last time
- Present relational data model: a way of representing data and constraints on that data
- SQL: a *declarative* language for creating, manipulating relational data
 - Declarative = say ***what*** you want, not ***how***

Motivation: why DBs and SQL?

- **Bigger data:** databases are optimized for processing data that is too large to fit in memory
- **More complex data:** CSV file is great for describing a single collection of entities. What about multiple types of entities? Relations between entities?
- **Common data processing patterns:**
 - Typical operations: *filter* rows/columns, *group* related rows, *aggregate* statistics, *join* different data sources
 - Codified into formal language (SQL)

Relational Database: Definitions

- *Relational database*: a set of *relations*
- *Relation*: made up of 2 parts:
 - *Schema* : name of relation, plus name and type/domain of each column.

Students(*sid*: string, *name*: string, *login*: string,
age: integer, *gpa*: real).

- *Instance* : the actual data at a given time

Relational instance: a table

Students

column,
attribute,
field

| sid | name | login | age | gpa |
|-------|-------|------------|-----|-----|
| 53666 | Jones | jones@cs | 18 | 3.4 |
| 53688 | Smith | smith@eecs | 18 | 3.2 |
| 53650 | Smith | smith@math | 19 | 3.8 |

row, tuple

Attribute value

- #rows = cardinality
- #fields = degree / arity

Synonyms

| Formal | Not-so-formal-1 | Not-so-formal 2 |
|-----------|-----------------|-----------------|
| Relation | Table | |
| Tuple | Row | Record |
| Attribute | Column | Field |
| Domain | Type | |

Example Database Instance

STUDENT

| sid | name |
|-----|------|
| 1 | Jill |
| 2 | Bo |
| 3 | Maya |

Takes

| sid | cid |
|-----|-----|
| 1 | 460 |
| 1 | 301 |
| 3 | 301 |

COURSE

| cid | title | sem |
|-----|-------|-----|
| 460 | DB | F16 |
| 480 | DS | S17 |
| 301 | OS | F16 |

PROFESSOR

| fid | name |
|-----|---------|
| 1 | Hay |
| 2 | Sommers |
| 8 | Mulry |

Teaches

| fid | cid |
|-----|-----|
| 1 | 460 |
| 1 | 480 |
| 2 | 301 |

Integrity constraints

- An **integrity constraint** (IC) is a condition that must be true for *any* instance of the DB
 - ICs specified when schema defined
 - ICs checked when relations are modified
- If the DBMS checks ICs, stored data is more faithful to real-world meaning (helps avoid data entry errors too!)

Keys

- **Keys** are a kind of integrity constraint
 - **Key**: field(s) such that no two distinct tuples can have same value
 - **Primary key**: if more than one key for relation, one key is chosen (by DBA) to be **primary key**
 - **Foreign key**: field(s) in one relation that is used to “refer” to tuple in another relation. (Must correspond to primary key of the second relation.) Like a “logical pointer.”
- Examples of key constraints in student courses database?

Example Database Instance

STUDENT

| sid | name |
|-----|------|
| 1 | Jill |
| 2 | Bo |
| 3 | Maya |

Takes

| sid | cid |
|-----|-----|
| 1 | 460 |
| 1 | 301 |
| 3 | 301 |

COURSE

| cid | title | sem |
|-----|-------|-----|
| 460 | DB | F14 |
| 480 | AI | S14 |
| 301 | OS | F14 |

PROFESSOR

| fid | name |
|-----|---------|
| 1 | Hay |
| 2 | Sommers |
| 8 | Mulry |

Teaches

| fid | cid |
|-----|-----|
| 1 | 460 |
| 1 | 480 |
| 2 | 301 |

What are the keys
of these relations?

Are there any
foreign keys?

Question: can a student apply to the same college twice?

CREATE TABLE

```
create table Student(  
  sID int,  
  sName text,  
  GPA real CHECK (gpa>0 AND gpa<=4.0),  
  sizeHS int,  
  PRIMARY KEY(sID)  
);
```

Integrity constraint:

GPA value must fall within this range

Key constraint: no two rows can have same key value

```
create table Apply(  
  sID int,  
  cName text,  
  major text,  
  decision text,  
  PRIMARY KEY (sID, cName, major),  
  FOREIGN KEY(sID) REFERENCES Student(sID),  
  FOREIGN KEY(cName) REFERENCES College(cName)  
);
```

Foreign key constraint: every sID value must match some student sID value found in Student; i.e., no “dangling” references

SQL Overview

See handout.