

Lecture 12: Linear Regression I & Gradient Descent

COSC 480 Data Science, Spring 2017
Michael Hay

Logistics

- Lab 4 out soon (late tonight)
- Textbook...
- Quiz tomorrow

Let's focus on modeling a *single* predictor variable




FIGURE 2.1. *The Advertising data set. The plot displays sales, in thousands of units, as a function of TV, radio, and newspaper budgets, in thousands of dollars, for 200 different markets. In each plot we show the simple least squares fit of sales to that variable, as described in Chapter 3. In other words, each blue line represents a simple model that can be used to predict sales using TV, radio, and newspaper, respectively.*




FIGURE 3.1. For the **Advertising** data, the least squares fit for the regression of **sales** onto **TV** is shown. The fit is found by minimizing the sum of squared errors. Each grey line segment represents an error, and the fit makes a compromise by averaging their squares. In this case a linear fit captures the essence of the relationship, although it is somewhat deficient in the left of the plot.

Algorithm 1 Gradient descent for simple linear regression

```
1: procedure GRADIENTDESCENT(cost function  $J$ , step size  $\eta$ , tolerance  $t$ )
2:   Initialize  $\beta_0, \beta_1$  arbitrarily.
3:   cost  $\leftarrow J(\beta_0, \beta_1)$ 
4:   repeat
5:      $\triangleright$  Use gradient to update coefficients:
6:      $\beta_0 \leftarrow \beta_0 - \eta \cdot \frac{\partial}{\partial \beta_0} J(\beta_0, \beta_1)$   $\triangleright$  Take step in opposite direction of gradient
7:      $\beta_1 \leftarrow \beta_1 - \eta \cdot \frac{\partial}{\partial \beta_1} J(\beta_0, \beta_1)$ 
8:     oldCost  $\leftarrow$  cost
9:     cost  $\leftarrow J(\beta_0, \beta_1)$ 
10:   until  $|\text{cost} - \text{oldCost}| < t$   $\triangleright$  Or some other “convergence” criterion
11:   Return  $\beta_0, \beta_1$ .
12: end procedure
```

Algorithm 2 Stochastic gradient descent for simple linear regression

```
1: procedure STOCHASTICGRADIENTDESCENT(dataset, step size  $\eta$ , tolerance  $t$ )
2:   Initialize  $\beta_0, \beta_1$  arbitrarily.
3:   cost  $\leftarrow J(\beta_0, \beta_1)$  ▷ This is computed using the dataset.
4:   repeat
5:     for  $i = 1 \dots n$  do
6:       Choose a random data point  $(x_i, y_i)$  from dataset. ▷ The “stochastic” part.
7:       ▷ Use gradient of squared error on point  $(x_i, y_i)$  to update coefficients:
8:        $\beta_0 \leftarrow \beta_0 + \eta \cdot 2(y_i - h_\beta(x_i))$ 
9:        $\beta_1 \leftarrow \beta_1 + \eta \cdot 2(y_i - h_\beta(x_i)) x_i$ 
10:    end for
11:    oldCost  $\leftarrow$  cost
12:    cost  $\leftarrow J(\beta_0, \beta_1)$ 
13:  until  $|\text{cost} - \text{oldCost}| < t$ 
14:  Return  $\beta_0, \beta_1$ .
15: end procedure
```
